

Practicing Hadoop File System (fs) Commands

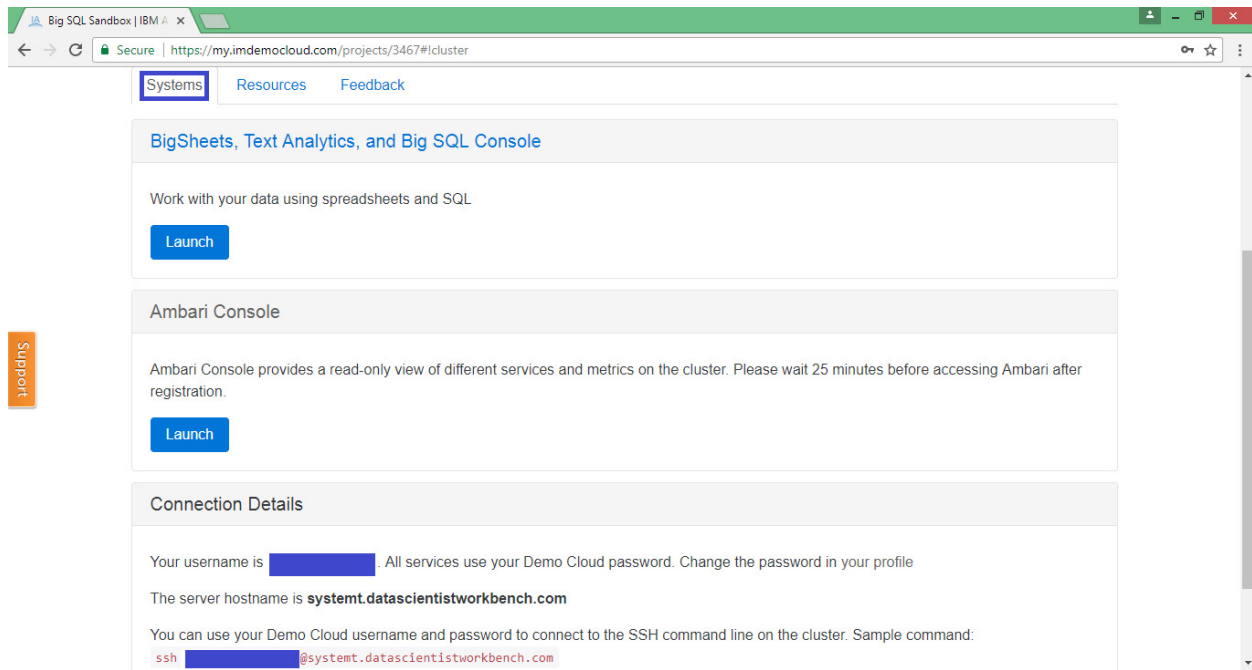
I generally learn and remember better, if I practice.

Putting together various Hadoop file system (fs) commands that I practiced in IBM Analytics Demo Cloud. If you want to sign up for IBM Analytics Demo Cloud, follow instructions in <https://my.imdemocloud.com>. Signing up was free when I start writing this article.

Adding few screenshots that will help to check the username needed to connect to IBM Analytics Demo Cloud.

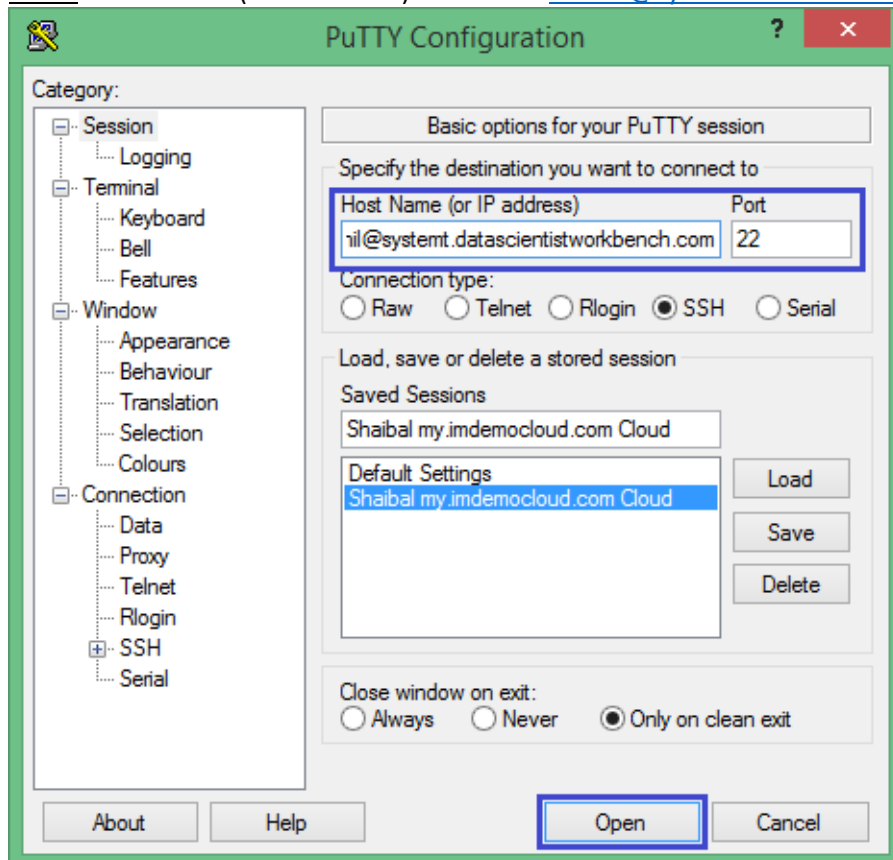
The first screenshot shows the IBM Analytics Demo Cloud login page. The browser address bar displays <https://my.imdemocloud.com>. The page header includes the IBM logo, 'Home', and 'Logout' links. A blue banner indicates 'Signed in successfully.' The main heading is 'IBM Analytics Demo Cloud' with the subtext 'Demo Cloud provides you with environments you can use to explore and experiment with new software capabilities.' A 'Support' link is visible on the left. Below the heading is the 'Your Projects' section, which contains a box labeled 'Big SQL Sandbox' with a blue arrow pointing to it.

The second screenshot shows the 'Big SQL Sandbox' page. The browser address bar displays <https://my.imdemocloud.com>. The page header includes the IBM logo, 'Home', 'Logout', and 'Welcome, SHAIBAL KANTI SHIL'. The main heading is 'Big SQL Sandbox' with the subtext 'Big SQL Sandbox is a large, shared environment for data science. You can use it to run R, SQL, Spark, and Hadoop jobs. It is a high performance cluster demonstrating the advantages of parallelized processing of big data sets.' A 'Support' link is visible on the left. Below the heading is a navigation bar with 'Systems', 'Resources', and 'Feedback' links. The 'Systems' link is highlighted with a blue box and a blue arrow pointing to it. Below the navigation bar is a section titled 'BigSheets, Text Analytics, and Big SQL Console' with the subtext 'Work with your data using spreadsheets and SQL' and a blue 'Launch' button.



I access my IBM Analytics Demo Cloud account from my windows 64-bit machine via PuTTY. In case you need PuTTY follow this URL: <https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>

NOTE: 'Host Name (or IP address)' looks like: [UserID@systemt.datascientistworkbench.com](https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html)

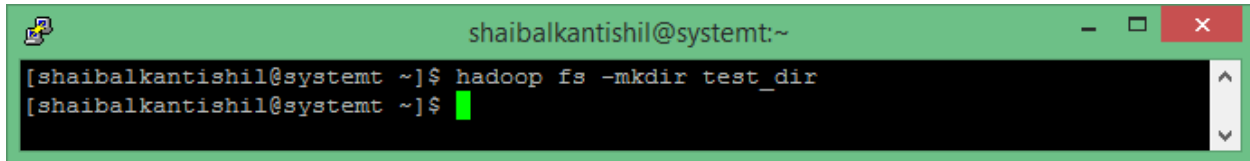


List of Hadoop File System (fs) commands along with description can be found at:

<https://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

I am adding screenshot after executing the commands.

mkdir

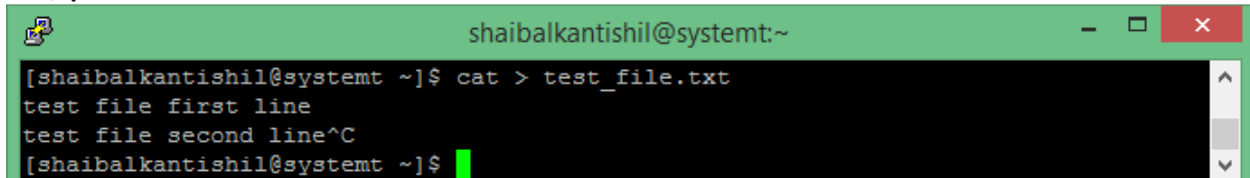


```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -mkdir test_dir
[shaibalkantishil@systemt ~]$

```

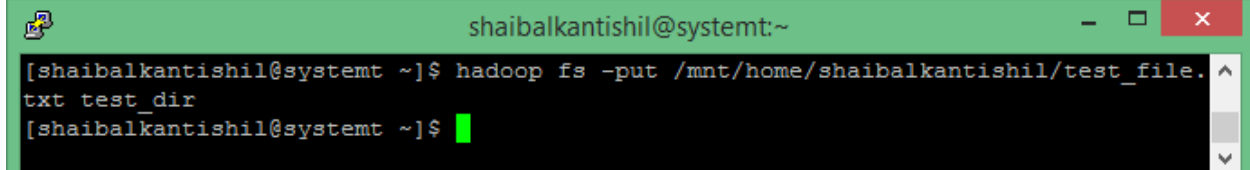
cat / put



```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ cat > test_file.txt
test file first line
test file second line^C
[shaibalkantishil@systemt ~]$

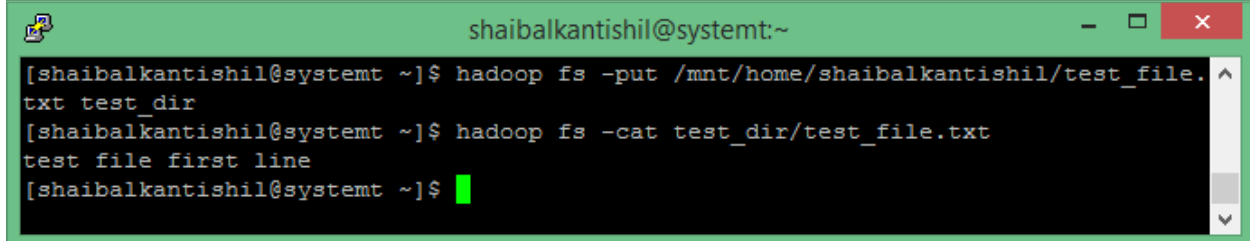
```



```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -put /mnt/home/shaibalkantishil/test_file.txt test_dir
[shaibalkantishil@systemt ~]$

```

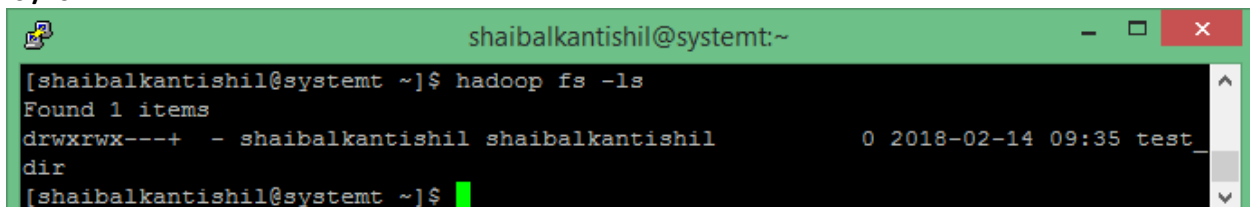


```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -put /mnt/home/shaibalkantishil/test_file.txt test_dir
[shaibalkantishil@systemt ~]$ hadoop fs -cat test_dir/test_file.txt
test file first line
[shaibalkantishil@systemt ~]$

```

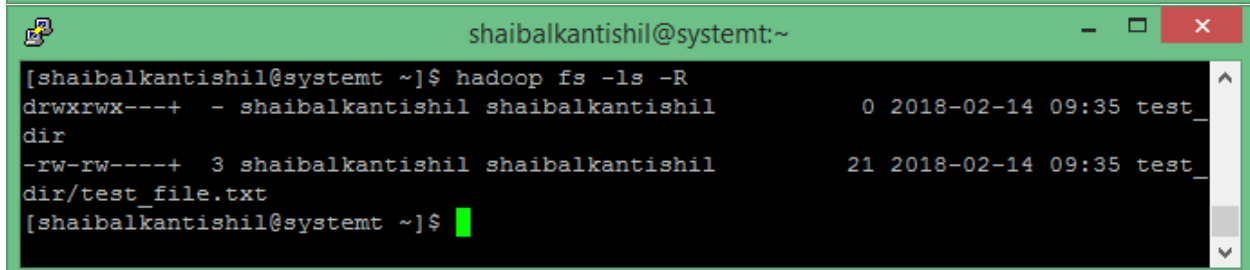
ls / lsr



```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -ls
Found 1 items
drwxrwx---+ - shaibalkantishil shaibalkantishil      0 2018-02-14 09:35 test_dir
[shaibalkantishil@systemt ~]$

```



```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -ls -R
drwxrwx---+ - shaibalkantishil shaibalkantishil      0 2018-02-14 09:35 test_dir
-rw-rw----+ 3 shaibalkantishil shaibalkantishil    21 2018-02-14 09:35 test_dir/test_file.txt
[shaibalkantishil@systemt ~]$

```

```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -lsr
lsr: DEPRECATED: Please use 'ls -R' instead.
drwxrwx---+ - shaibalkantishil shaibalkantishil      0 2018-02-14 09:35 test_
dir
-rw-rw----+ 3 shaibalkantishil shaibalkantishil    21 2018-02-14 09:35 test_
dir/test_file.txt
[shaibalkantishil@systemt ~]$

```

du / dus

```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -du test_dir/test_file.txt
21 test_dir/test_file.txt
[shaibalkantishil@systemt ~]$ hadoop fs -dus test_dir/test_file.txt
dus: DEPRECATED: Please use 'du -s' instead.
21 test_dir/test_file.txt
[shaibalkantishil@systemt ~]$ hadoop fs -du -s test_dir/test_file.txt
21 test_dir/test_file.txt
[shaibalkantishil@systemt ~]$

```

cp

```

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -cp test_dir/test_file.txt test_dir/test_
_copied_file.txt
18/02/23 15:06:18 WARN hdfs.DFSClient: DFSInputStream has been closed already
[shaibalkantishil@systemt ~]$

```

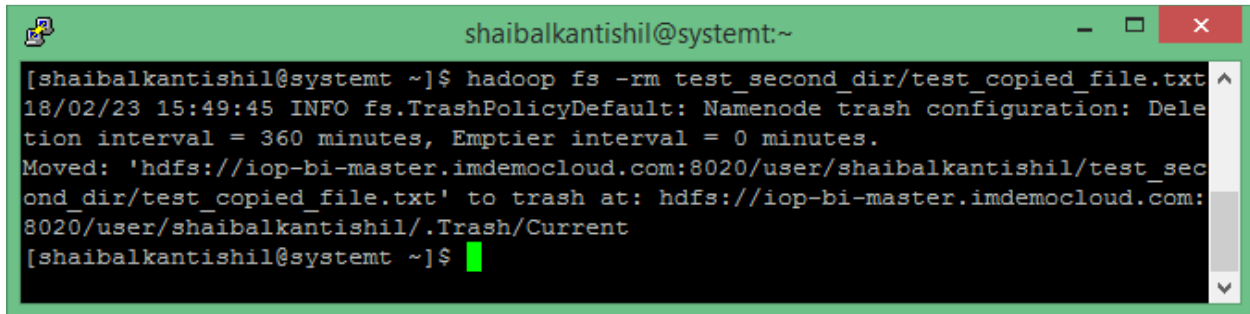
mv

```

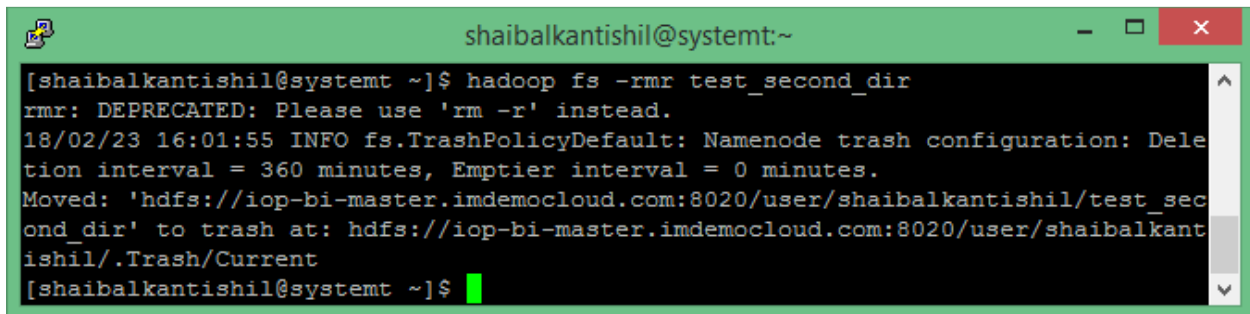
shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -mkdir test_second_dir

shaibalkantishil@systemt:~
[shaibalkantishil@systemt ~]$ hadoop fs -mv test_dir/test_copied_file.txt test_s
econd_dir
[shaibalkantishil@systemt ~]$

```

rm

```
shaibalkantishil@systemt:~  
[shaibalkantishil@systemt ~]$ hadoop fs -rm test_second_dir/test_copied_file.txt  
18/02/23 15:49:45 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.  
Moved: 'hdfs://iop-bi-master.imdemocloud.com:8020/user/shaibalkantishil/test_second_dir/test_copied_file.txt' to trash at: hdfs://iop-bi-master.imdemocloud.com:8020/user/shaibalkantishil/.Trash/Current  
[shaibalkantishil@systemt ~]$
```

rmr

```
shaibalkantishil@systemt:~  
[shaibalkantishil@systemt ~]$ hadoop fs -rmr test_second_dir  
rmr: DEPRECATED: Please use 'rm -r' instead.  
18/02/23 16:01:55 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.  
Moved: 'hdfs://iop-bi-master.imdemocloud.com:8020/user/shaibalkantishil/test_second_dir' to trash at: hdfs://iop-bi-master.imdemocloud.com:8020/user/shaibalkantishil/.Trash/Current  
[shaibalkantishil@systemt ~]$
```